

dog_app

April 19, 2020

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/.*"))
        dog_files = np.array(glob("/data/dog_images/*/.*"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [3]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [4]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell)

```
In [5]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

##-## Do NOT modify the code above this line. ##-##

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.

count_humans = 0
count_dogs = 0

for file in human_files_short:
    if face_detector(file) == True:
        count_humans +=1

for file in dog_files_short:
    if face_detector(file) == True:
        count_dogs +=1

print('%.1f%% images of the first 100 human_files were detected as human face.' % count_humans)
print('%.1f%% images of the first 100 dog_files were detected as human face.' % count_dogs)

98.0% images of the first 100 human_files were detected as human face.
17.0% images of the first 100 dog_files were detected as human face.
```

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [6]: ### (Optional)  
       ### TODO: Test performance of another face detection algorithm.  
       ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [7]: import torch  
       import torchvision.models as models  
  
       # define VGG16 model  
       VGG16 = models.vgg16(pretrained=True)  
  
       # check if CUDA is available  
       use_cuda = torch.cuda.is_available()  
  
       # move model to GPU if CUDA is available  
       if use_cuda:  
           VGG16 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth  
100%|| 553433881/553433881 [00:05<00:00, 98324211.51it/s]
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```

In [8]: from PIL import Image
import torchvision.transforms as transforms

def load_image(img_path):
    image = Image.open(img_path).convert('RGB')
    # resize the input image into 224x224 because VGG16 takes only 224x224 pixel image
    in_transform = transforms.Compose([transforms.Resize(size=(224,224)),
                                      transforms.ToTensor()
                                      ])
    image = in_transform(image)[:3,:,:].unsqueeze(0)
    return image

In [9]: from PIL import Image
import torchvision.transforms as transforms

from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

def VGG16_predict(img_path):
    '''
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    '''

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image
    img = load_image(img_path)
    if use_cuda:
        img = img.cuda()
    ret = VGG16(img)

    return torch.max(ret,1)[1].item() # predicted class index

In [10]: VGG16_predict(dog_files_short[0])

Out[10]: 243

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is

predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

```
In [11]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    idx = VGG16_predict(img_path)
    return idx >= 151 and idx <= 268
```

```
In [12]: print(dog_detector(dog_files_short[0]))
print(dog_detector(human_files_short[0]))
```

```
True
False
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer: 0% of dogs in `human_files_short` and 93% of dogs in `dog_files_short`

```
In [13]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.
```

```
def dog_detector_test(files):
    detection_cnt = 0;
    total_cnt = len(files)
    for file in files:
        detection_cnt += dog_detector(file)
    return detection_cnt, total_cnt
```

```
In [14]: print("detect a dog in human_files: {} / {}".format(dog_detector_test(human_files_short)
print("detect a dog in dog_files: {} / {}".format(dog_detector_test(dog_files_short)[0]
```

```
detect a dog in human_files: 0 / 100
detect a dog in dog_files: 93 / 100
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [15]: ### (Optional)
### TODO: Report the performance of another pre-trained network.
### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [16]: import os
         from torchvision import datasets

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes
         ImageFile.LOAD_TRUNCATED_IMAGES = True
```



```

batch_size = 20
num_workers = 0 # default is 0

data_dir = '/data/dog_images/'
train_dir = os.path.join(data_dir, 'train')
valid_dir = os.path.join(data_dir, 'valid')
test_dir = os.path.join(data_dir, 'test')

standard_normalization = transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.229, 0.229])

data_transforms = {'train': transforms.Compose([transforms.RandomResizedCrop(224),
                                                transforms.RandomHorizontalFlip(),
                                                transforms.ToTensor(),
                                                standard_normalization
                                                ]),
                   'val': transforms.Compose([transforms.Resize(256),
                                                transforms.CenterCrop(224),
                                                transforms.ToTensor(),
                                                standard_normalization]),
                   'test': transforms.Compose([transforms.Resize(size=(224,224)),
                                                transforms.ToTensor(),
                                                standard_normalization])
                   }

In [17]: train_data = datasets.ImageFolder(train_dir, transform=data_transforms['train'])
        valid_data = datasets.ImageFolder(valid_dir, transform=data_transforms['val'])
        test_data = datasets.ImageFolder(test_dir, transform = data_transforms['test'])

train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, num_workers=num_workers)
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size, num_workers=num_workers)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, num_workers=num_workers)

loader_scratch = {'train': train_loader, 'valid': valid_loader, 'test': test_loader}

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer: As resizing the image data is one of the important part of preprocessing because different images have different pixel sizes, so i resize my images dataset using RandomResizedCrop and i set the resize value to be 224. Augmenting the data also helps when we have imbalanced dataset problem. So i applied RandomHorizontalClip function to our training data only because we will train our data on that train dataset only rest for valid and test dataset we will not require to do augmentation.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```
In [18]: num_classes = 133

In [19]: import torch.nn as nn
import torch.nn.functional as F
import numpy as np

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1 = nn.Conv2d(3, 32, 3, stride=2, padding=1)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=2, padding=1)
        self.conv3 = nn.Conv2d(64, 128, 3, padding=1)

        # pool
        self.pool = nn.MaxPool2d(2, 2)

        # fully-connected
        self.fc1 = nn.Linear(7*7*128, 500)
        self.fc2 = nn.Linear(500, num_classes)

        # drop-out
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        ## Define forward behavior
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x = F.relu(self.conv2(x))
        x = self.pool(x)
        x = F.relu(self.conv3(x))
        x = self.pool(x)
        (_, C, H, W) = x.data.size()

        # flatten
        x = x.view(-1, C * H * W)

        x = self.dropout(x)
        x = F.relu(self.fc1(x))

        x = self.dropout(x)
        x = self.fc2(x)
        return x
```

```

    ##-## You so NOT have to modify the code below this line. ##-##

    # instantiate the CNN
    model_scratch = Net()
    print(model_scratch)
    # move tensors to GPU if CUDA is available
    if use_cuda:
        model_scratch.cuda()

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=6272, out_features=500, bias=True)
  (fc2): Linear(in_features=500, out_features=133, bias=True)
  (dropout): Dropout(p=0.3)
)

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer: I have used 3 conv2d layer for architecture with kernel size = 3, stride = 2, padding = 1. The first 2 layer implemented with a kernel size of 3 and stride 2 and padding of 1. The 3rd layer is implemented with kernel size of 3 and padding of 1 but i have not implemented stride operation on 3rd layer. After each layer i implemented the maxpooling layer and at last i have used dropout regularization with value 0.3. Dropout prevents our algorithm from overfitting. I have used RELU activation function at the hidden layers of our CNN architecture. RELU activation function gives values between 0 and 1. I have used 2 fully connected layers. Last fully connected layer gives the predictions from one of the 133 classes.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [20]: import torch.optim as optim

    ### TODO: select loss function
    criterion_scratch = nn.CrossEntropyLoss()

    ### TODO: select optimizer
    optimizer_scratch = optim.Adam(model_scratch.parameters(), lr=0.0001)

```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [21]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path, last_val
        """returns trained model"""
        # initialize tracker for minimum validation loss
        if last_validation_loss is not None:
            valid_loss_min = last_validation_loss
        else:
            valid_loss_min = np.Inf

        for epoch in range(1, n_epochs+1):
            # initialize variables to monitor training and validation loss
            train_loss = 0.0
            valid_loss = 0.0

            #####
            # train the model #
            #####
            model.train()
            for batch_idx, (data, target) in enumerate(loaders['train']):
                # move to GPU
                if use_cuda:
                    data, target = data.cuda(), target.cuda()
                ## find the loss and update the model parameters accordingly
                ## record the average training loss, using something like
                ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

                # initialize weights to zero
                optimizer.zero_grad()

                output = model(data)

                # calculate loss
                loss = criterion(output, target)

                # back prop
                loss.backward()

                # grad
                optimizer.step()

                train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

            if batch_idx % 100 == 0:
                print('Epoch %d, Batch %d loss: %.6f' %
                      (epoch, batch_idx + 1, train_loss))

            #####
            # validate the model #
            #####

```

```

model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss
    output = model(data)
    loss = criterion(output, target)
    valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss < valid_loss_min:
    torch.save(model.state_dict(), save_path)
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
        valid_loss_min,
        valid_loss))
    valid_loss_min = valid_loss

# return trained model
return model

```

```
In [22]: model_scratch = train(40, loader_scratch, model_scratch, optimizer_scratch, criterion_
```

```

Epoch 1, Batch 1 loss: 4.871049
Epoch 1, Batch 101 loss: 4.893119
Epoch 1, Batch 201 loss: 4.885406
Epoch 1, Batch 301 loss: 4.867074
Epoch: 1          Training Loss: 4.864035          Validation Loss: 4.786391
Validation loss decreased (inf --> 4.786391). Saving model ...
Epoch 2, Batch 1 loss: 4.781516
Epoch 2, Batch 101 loss: 4.751566
Epoch 2, Batch 201 loss: 4.721057
Epoch 2, Batch 301 loss: 4.705526
Epoch: 2          Training Loss: 4.704655          Validation Loss: 4.559371
Validation loss decreased (4.786391 --> 4.559371). Saving model ...
Epoch 3, Batch 1 loss: 4.849067
Epoch 3, Batch 101 loss: 4.630446
Epoch 3, Batch 201 loss: 4.603781
Epoch 3, Batch 301 loss: 4.596972
Epoch: 3          Training Loss: 4.597475          Validation Loss: 4.470085

```

Validation loss decreased (4.559371 --> 4.470085). Saving model ...
 Epoch 4, Batch 1 loss: 4.419483
 Epoch 4, Batch 101 loss: 4.535429
 Epoch 4, Batch 201 loss: 4.543274
 Epoch 4, Batch 301 loss: 4.535500
 Epoch: 4 Training Loss: 4.531193 Validation Loss: 4.384349
 Validation loss decreased (4.470085 --> 4.384349). Saving model ...
 Epoch 5, Batch 1 loss: 4.461885
 Epoch 5, Batch 101 loss: 4.467141
 Epoch 5, Batch 201 loss: 4.457199
 Epoch 5, Batch 301 loss: 4.452735
 Epoch: 5 Training Loss: 4.450023 Validation Loss: 4.286588
 Validation loss decreased (4.384349 --> 4.286588). Saving model ...
 Epoch 6, Batch 1 loss: 4.386744
 Epoch 6, Batch 101 loss: 4.389340
 Epoch 6, Batch 201 loss: 4.402273
 Epoch 6, Batch 301 loss: 4.414482
 Epoch: 6 Training Loss: 4.413668 Validation Loss: 4.272698
 Validation loss decreased (4.286588 --> 4.272698). Saving model ...
 Epoch 7, Batch 1 loss: 4.475667
 Epoch 7, Batch 101 loss: 4.334605
 Epoch 7, Batch 201 loss: 4.358491
 Epoch 7, Batch 301 loss: 4.357951
 Epoch: 7 Training Loss: 4.356931 Validation Loss: 4.196947
 Validation loss decreased (4.272698 --> 4.196947). Saving model ...
 Epoch 8, Batch 1 loss: 4.300153
 Epoch 8, Batch 101 loss: 4.298705
 Epoch 8, Batch 201 loss: 4.298914
 Epoch 8, Batch 301 loss: 4.295869
 Epoch: 8 Training Loss: 4.297266 Validation Loss: 4.116701
 Validation loss decreased (4.196947 --> 4.116701). Saving model ...
 Epoch 9, Batch 1 loss: 4.138509
 Epoch 9, Batch 101 loss: 4.234140
 Epoch 9, Batch 201 loss: 4.247304
 Epoch 9, Batch 301 loss: 4.245508
 Epoch: 9 Training Loss: 4.244715 Validation Loss: 4.052224
 Validation loss decreased (4.116701 --> 4.052224). Saving model ...
 Epoch 10, Batch 1 loss: 4.305787
 Epoch 10, Batch 101 loss: 4.179132
 Epoch 10, Batch 201 loss: 4.197581
 Epoch 10, Batch 301 loss: 4.203821
 Epoch: 10 Training Loss: 4.202878 Validation Loss: 4.025463
 Validation loss decreased (4.052224 --> 4.025463). Saving model ...
 Epoch 11, Batch 1 loss: 3.869030
 Epoch 11, Batch 101 loss: 4.157185
 Epoch 11, Batch 201 loss: 4.169479
 Epoch 11, Batch 301 loss: 4.175796
 Epoch: 11 Training Loss: 4.169549 Validation Loss: 3.970160

Validation loss decreased (4.025463 --> 3.970160). Saving model ...
 Epoch 12, Batch 1 loss: 3.899272
 Epoch 12, Batch 101 loss: 4.119015
 Epoch 12, Batch 201 loss: 4.103691
 Epoch 12, Batch 301 loss: 4.101757
 Epoch: 12 Training Loss: 4.104415 Validation Loss: 3.932025
 Validation loss decreased (3.970160 --> 3.932025). Saving model ...
 Epoch 13, Batch 1 loss: 3.868602
 Epoch 13, Batch 101 loss: 4.068631
 Epoch 13, Batch 201 loss: 4.070665
 Epoch 13, Batch 301 loss: 4.061615
 Epoch: 13 Training Loss: 4.063232 Validation Loss: 3.899366
 Validation loss decreased (3.932025 --> 3.899366). Saving model ...
 Epoch 14, Batch 1 loss: 4.004123
 Epoch 14, Batch 101 loss: 4.018809
 Epoch 14, Batch 201 loss: 4.028363
 Epoch 14, Batch 301 loss: 4.042857
 Epoch: 14 Training Loss: 4.038366 Validation Loss: 3.867612
 Validation loss decreased (3.899366 --> 3.867612). Saving model ...
 Epoch 15, Batch 1 loss: 3.664016
 Epoch 15, Batch 101 loss: 3.969246
 Epoch 15, Batch 201 loss: 3.970093
 Epoch 15, Batch 301 loss: 3.990197
 Epoch: 15 Training Loss: 3.990407 Validation Loss: 3.834262
 Validation loss decreased (3.867612 --> 3.834262). Saving model ...
 Epoch 16, Batch 1 loss: 3.723940
 Epoch 16, Batch 101 loss: 3.971588
 Epoch 16, Batch 201 loss: 3.976279
 Epoch 16, Batch 301 loss: 3.962242
 Epoch: 16 Training Loss: 3.968453 Validation Loss: 3.805078
 Validation loss decreased (3.834262 --> 3.805078). Saving model ...
 Epoch 17, Batch 1 loss: 3.696424
 Epoch 17, Batch 101 loss: 3.919991
 Epoch 17, Batch 201 loss: 3.914757
 Epoch 17, Batch 301 loss: 3.931229
 Epoch: 17 Training Loss: 3.927326 Validation Loss: 3.753586
 Validation loss decreased (3.805078 --> 3.753586). Saving model ...
 Epoch 18, Batch 1 loss: 3.783653
 Epoch 18, Batch 101 loss: 3.846440
 Epoch 18, Batch 201 loss: 3.889733
 Epoch 18, Batch 301 loss: 3.898976
 Epoch: 18 Training Loss: 3.905682 Validation Loss: 3.753460
 Validation loss decreased (3.753586 --> 3.753460). Saving model ...
 Epoch 19, Batch 1 loss: 3.314863
 Epoch 19, Batch 101 loss: 3.896262
 Epoch 19, Batch 201 loss: 3.902553
 Epoch 19, Batch 301 loss: 3.881989
 Epoch: 19 Training Loss: 3.875132 Validation Loss: 3.711045

Validation loss decreased (3.753460 --> 3.711045). Saving model ...
 Epoch 20, Batch 1 loss: 3.673454
 Epoch 20, Batch 101 loss: 3.822494
 Epoch 20, Batch 201 loss: 3.841267
 Epoch 20, Batch 301 loss: 3.837610
 Epoch: 20 Training Loss: 3.844916 Validation Loss: 3.709469
 Validation loss decreased (3.711045 --> 3.709469). Saving model ...
 Epoch 21, Batch 1 loss: 3.816968
 Epoch 21, Batch 101 loss: 3.781257
 Epoch 21, Batch 201 loss: 3.804330
 Epoch 21, Batch 301 loss: 3.802641
 Epoch: 21 Training Loss: 3.805174 Validation Loss: 3.669756
 Validation loss decreased (3.709469 --> 3.669756). Saving model ...
 Epoch 22, Batch 1 loss: 3.867445
 Epoch 22, Batch 101 loss: 3.813908
 Epoch 22, Batch 201 loss: 3.803663
 Epoch 22, Batch 301 loss: 3.786221
 Epoch: 22 Training Loss: 3.785575 Validation Loss: 3.676198
 Epoch 23, Batch 1 loss: 3.669495
 Epoch 23, Batch 101 loss: 3.757560
 Epoch 23, Batch 201 loss: 3.744747
 Epoch 23, Batch 301 loss: 3.748357
 Epoch: 23 Training Loss: 3.749282 Validation Loss: 3.664670
 Validation loss decreased (3.669756 --> 3.664670). Saving model ...
 Epoch 24, Batch 1 loss: 3.860677
 Epoch 24, Batch 101 loss: 3.674917
 Epoch 24, Batch 201 loss: 3.699544
 Epoch 24, Batch 301 loss: 3.726098
 Epoch: 24 Training Loss: 3.718968 Validation Loss: 3.636617
 Validation loss decreased (3.664670 --> 3.636617). Saving model ...
 Epoch 25, Batch 1 loss: 3.915054
 Epoch 25, Batch 101 loss: 3.700215
 Epoch 25, Batch 201 loss: 3.732625
 Epoch 25, Batch 301 loss: 3.728964
 Epoch: 25 Training Loss: 3.727643 Validation Loss: 3.599848
 Validation loss decreased (3.636617 --> 3.599848). Saving model ...
 Epoch 26, Batch 1 loss: 3.663528
 Epoch 26, Batch 101 loss: 3.650700
 Epoch 26, Batch 201 loss: 3.672613
 Epoch 26, Batch 301 loss: 3.660950
 Epoch: 26 Training Loss: 3.665729 Validation Loss: 3.588993
 Validation loss decreased (3.599848 --> 3.588993). Saving model ...
 Epoch 27, Batch 1 loss: 4.046014
 Epoch 27, Batch 101 loss: 3.658147
 Epoch 27, Batch 201 loss: 3.658832
 Epoch 27, Batch 301 loss: 3.632631
 Epoch: 27 Training Loss: 3.648492 Validation Loss: 3.617462
 Epoch 28, Batch 1 loss: 3.133695

Epoch 28, Batch 101 loss: 3.617220
 Epoch 28, Batch 201 loss: 3.628532
 Epoch 28, Batch 301 loss: 3.626976
 Epoch: 28 Training Loss: 3.627260 Validation Loss: 3.578496
 Validation loss decreased (3.588993 --> 3.578496). Saving model ...
 Epoch 29, Batch 1 loss: 3.431058
 Epoch 29, Batch 101 loss: 3.524776
 Epoch 29, Batch 201 loss: 3.563466
 Epoch 29, Batch 301 loss: 3.565335
 Epoch: 29 Training Loss: 3.567729 Validation Loss: 3.627954
 Epoch 30, Batch 1 loss: 3.966942
 Epoch 30, Batch 101 loss: 3.619689
 Epoch 30, Batch 201 loss: 3.611555
 Epoch 30, Batch 301 loss: 3.592018
 Epoch: 30 Training Loss: 3.603190 Validation Loss: 3.556231
 Validation loss decreased (3.578496 --> 3.556231). Saving model ...
 Epoch 31, Batch 1 loss: 3.824828
 Epoch 31, Batch 101 loss: 3.538270
 Epoch 31, Batch 201 loss: 3.564185
 Epoch 31, Batch 301 loss: 3.559955
 Epoch: 31 Training Loss: 3.562279 Validation Loss: 3.517561
 Validation loss decreased (3.556231 --> 3.517561). Saving model ...
 Epoch 32, Batch 1 loss: 3.483950
 Epoch 32, Batch 101 loss: 3.516318
 Epoch 32, Batch 201 loss: 3.515965
 Epoch 32, Batch 301 loss: 3.521655
 Epoch: 32 Training Loss: 3.534929 Validation Loss: 3.519052
 Epoch 33, Batch 1 loss: 3.179828
 Epoch 33, Batch 101 loss: 3.532397
 Epoch 33, Batch 201 loss: 3.533552
 Epoch 33, Batch 301 loss: 3.532106
 Epoch: 33 Training Loss: 3.532175 Validation Loss: 3.486520
 Validation loss decreased (3.517561 --> 3.486520). Saving model ...
 Epoch 34, Batch 1 loss: 3.470614
 Epoch 34, Batch 101 loss: 3.464051
 Epoch 34, Batch 201 loss: 3.472976
 Epoch 34, Batch 301 loss: 3.476699
 Epoch: 34 Training Loss: 3.488072 Validation Loss: 3.505313
 Epoch 35, Batch 1 loss: 3.099227
 Epoch 35, Batch 101 loss: 3.453094
 Epoch 35, Batch 201 loss: 3.491020
 Epoch 35, Batch 301 loss: 3.491463
 Epoch: 35 Training Loss: 3.489741 Validation Loss: 3.476178
 Validation loss decreased (3.486520 --> 3.476178). Saving model ...
 Epoch 36, Batch 1 loss: 3.070490
 Epoch 36, Batch 101 loss: 3.439601
 Epoch 36, Batch 201 loss: 3.468182
 Epoch 36, Batch 301 loss: 3.465155

```

Epoch: 36          Training Loss: 3.462722          Validation Loss: 3.455202
Validation loss decreased (3.476178 --> 3.455202). Saving model ...
Epoch 37, Batch 1 loss: 3.788702
Epoch 37, Batch 101 loss: 3.420576
Epoch 37, Batch 201 loss: 3.460891
Epoch 37, Batch 301 loss: 3.454637
Epoch: 37          Training Loss: 3.456223          Validation Loss: 3.464204
Epoch 38, Batch 1 loss: 3.067639
Epoch 38, Batch 101 loss: 3.419470
Epoch 38, Batch 201 loss: 3.436675
Epoch 38, Batch 301 loss: 3.440780
Epoch: 38          Training Loss: 3.448012          Validation Loss: 3.466990
Epoch 39, Batch 1 loss: 3.400626
Epoch 39, Batch 101 loss: 3.445652
Epoch 39, Batch 201 loss: 3.425551
Epoch 39, Batch 301 loss: 3.432493
Epoch: 39          Training Loss: 3.437361          Validation Loss: 3.438279
Validation loss decreased (3.455202 --> 3.438279). Saving model ...
Epoch 40, Batch 1 loss: 3.243506
Epoch 40, Batch 101 loss: 3.434815
Epoch 40, Batch 201 loss: 3.412618
Epoch 40, Batch 301 loss: 3.404395
Epoch: 40          Training Loss: 3.410262          Validation Loss: 3.423297
Validation loss decreased (3.438279 --> 3.423297). Saving model ...

```

```

In [23]: # load the model that got the best validation accuracy
         model_scratch.load_state_dict(torch.load('model_scratches.pt'))

```

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```

In [24]: def test(loaders, model, criterion, use_cuda):

         # monitor test loss and accuracy
         test_loss = 0.
         correct = 0.
         total = 0.

         model.eval()
         for batch_idx, (data, target) in enumerate(loaders['test']):
             # move to GPU
             if use_cuda:
                 data, target = data.cuda(), target.cuda()
             # forward pass: compute predicted outputs by passing inputs to the model
             output = model(data)

```

```

        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    test(loader_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.666922

Test Accuracy: 16% (138/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [25]: ## TODO: Specify data loaders
         loader_transfer = loader_scratch.copy()

```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [26]: import torchvision.models as models
         import torch.nn as nn

```

```

## TODO: Specify model architecture
model_transfer = models.resnet50(pretrained=True)

for param in model_transfer.parameters():
    param.requires_grad = False

model_transfer.fc = nn.Linear(2048,133,bias=True)

fc_parameters = model_transfer.fc.parameters()

for param in fc_parameters:
    param.requires_grad = True

model_transfer

if use_cuda:
    model_transfer = model_transfer.cuda()

```

Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/100%|| 102502400/102502400 [00:01<00:00, 85765821.99it/s]

In [27]: model_transfer

```

Out[27]: ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

        (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (layer2): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=True)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(layer3): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=True)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
  (1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=True)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=True)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=True)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
)

```

```

)
(4): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
(fc): Linear(in_features=2048, out_features=133, bias=True)
)

```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: After trying using VGG16 pretrained model for our custom CNN i get less accuracy but when i used Transfer Learning with resnet50 pretrained model even after less epochs i get great accuracy as compare to our custom CNN. In our custom CNN even after 40 epochs i get accuracy of 16% and when i used transfer learning even with epoch=20 i get accuracy of 85%. reset also prevents overfitting problems so i think it is the best possible solution to our problem

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```

In [28]: criterion_transfer = nn.CrossEntropyLoss()
        optimizer_transfer = optim.Adam(model_transfer.fc.parameters(), lr = 0.0001)

```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```

In [29]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""

        valid_loss_min = np.Inf

        for epoch in range(1, n_epochs+1):

            train_loss = 0.0
            valid_loss = 0.0

            model.train()
            for batch_idx, (data, target) in enumerate(loaders['train']):

                if use_cuda:
                    data, target = data.cuda(), target.cuda()

                optimizer.zero_grad()

```



```

        output = model(data)

        loss = criterion(output, target)

        loss.backward()

        optimizer.step()

        train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        if batch_idx % 100 == 0:
            print('Epoch %d, Batch %d loss: %.6f' %
                  (epoch, batch_idx + 1, train_loss))

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['valid']):

        if use_cuda:
            data, target = data.cuda(), target.cuda()

        output = model(data)
        loss = criterion(output, target)
        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    if valid_loss < valid_loss_min:
        torch.save(model.state_dict(), save_path)
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
            valid_loss_min,
            valid_loss))
        valid_loss_min = valid_loss

    # return trained model
    return model

```

```

In [30]: # train the model
         model_transfer = train(20, loader_transfer, model_transfer, optimizer_transfer, criteri

         # load the model that got the best validation accuracy (uncomment the line below)
         #model_transfer.load_state_dict(torch.load('model_transfer.pt'))

Epoch 1, Batch 1 loss: 4.946682
Epoch 1, Batch 101 loss: 4.721754
Epoch 1, Batch 201 loss: 4.472858
Epoch 1, Batch 301 loss: 4.254427
Epoch: 1          Training Loss: 4.191656          Validation Loss: 3.134689
Validation loss decreased (inf --> 3.134689).  Saving model ...
Epoch 2, Batch 1 loss: 3.408025
Epoch 2, Batch 101 loss: 3.316442
Epoch 2, Batch 201 loss: 3.180854
Epoch 2, Batch 301 loss: 3.051431
Epoch: 2          Training Loss: 3.014789          Validation Loss: 2.058029
Validation loss decreased (3.134689 --> 2.058029).  Saving model ...
Epoch 3, Batch 1 loss: 2.375204
Epoch 3, Batch 101 loss: 2.491588
Epoch 3, Batch 201 loss: 2.393599
Epoch 3, Batch 301 loss: 2.319348
Epoch: 3          Training Loss: 2.296501          Validation Loss: 1.485861
Validation loss decreased (2.058029 --> 1.485861).  Saving model ...
Epoch 4, Batch 1 loss: 2.178491
Epoch 4, Batch 101 loss: 1.930932
Epoch 4, Batch 201 loss: 1.916455
Epoch 4, Batch 301 loss: 1.891082
Epoch: 4          Training Loss: 1.876997          Validation Loss: 1.141196
Validation loss decreased (1.485861 --> 1.141196).  Saving model ...
Epoch 5, Batch 1 loss: 1.679326
Epoch 5, Batch 101 loss: 1.684159
Epoch 5, Batch 201 loss: 1.654349
Epoch 5, Batch 301 loss: 1.634781
Epoch: 5          Training Loss: 1.625985          Validation Loss: 0.957861
Validation loss decreased (1.141196 --> 0.957861).  Saving model ...
Epoch 6, Batch 1 loss: 1.200581
Epoch 6, Batch 101 loss: 1.489776
Epoch 6, Batch 201 loss: 1.475439
Epoch 6, Batch 301 loss: 1.459616
Epoch: 6          Training Loss: 1.446528          Validation Loss: 0.834507
Validation loss decreased (0.957861 --> 0.834507).  Saving model ...
Epoch 7, Batch 1 loss: 1.191185
Epoch 7, Batch 101 loss: 1.313285
Epoch 7, Batch 201 loss: 1.326967
Epoch 7, Batch 301 loss: 1.318720
Epoch: 7          Training Loss: 1.313999          Validation Loss: 0.749984
Validation loss decreased (0.834507 --> 0.749984).  Saving model ...

```

Epoch 8, Batch 1 loss: 1.310118
 Epoch 8, Batch 101 loss: 1.295108
 Epoch 8, Batch 201 loss: 1.261436
 Epoch 8, Batch 301 loss: 1.236604
 Epoch: 8 Training Loss: 1.230660 Validation Loss: 0.688011
 Validation loss decreased (0.749984 --> 0.688011). Saving model ...
 Epoch 9, Batch 1 loss: 1.020746
 Epoch 9, Batch 101 loss: 1.138199
 Epoch 9, Batch 201 loss: 1.144033
 Epoch 9, Batch 301 loss: 1.144463
 Epoch: 9 Training Loss: 1.142856 Validation Loss: 0.639458
 Validation loss decreased (0.688011 --> 0.639458). Saving model ...
 Epoch 10, Batch 1 loss: 0.947936
 Epoch 10, Batch 101 loss: 1.100933
 Epoch 10, Batch 201 loss: 1.108604
 Epoch 10, Batch 301 loss: 1.102136
 Epoch: 10 Training Loss: 1.102283 Validation Loss: 0.591206
 Validation loss decreased (0.639458 --> 0.591206). Saving model ...
 Epoch 11, Batch 1 loss: 1.030619
 Epoch 11, Batch 101 loss: 1.039137
 Epoch 11, Batch 201 loss: 1.043089
 Epoch 11, Batch 301 loss: 1.037969
 Epoch: 11 Training Loss: 1.035946 Validation Loss: 0.554545
 Validation loss decreased (0.591206 --> 0.554545). Saving model ...
 Epoch 12, Batch 1 loss: 1.012159
 Epoch 12, Batch 101 loss: 0.941199
 Epoch 12, Batch 201 loss: 0.980108
 Epoch 12, Batch 301 loss: 0.987987
 Epoch: 12 Training Loss: 0.991515 Validation Loss: 0.533829
 Validation loss decreased (0.554545 --> 0.533829). Saving model ...
 Epoch 13, Batch 1 loss: 1.028037
 Epoch 13, Batch 101 loss: 0.953201
 Epoch 13, Batch 201 loss: 0.954135
 Epoch 13, Batch 301 loss: 0.957232
 Epoch: 13 Training Loss: 0.958578 Validation Loss: 0.525996
 Validation loss decreased (0.533829 --> 0.525996). Saving model ...
 Epoch 14, Batch 1 loss: 1.021285
 Epoch 14, Batch 101 loss: 0.971218
 Epoch 14, Batch 201 loss: 0.954885
 Epoch 14, Batch 301 loss: 0.935948
 Epoch: 14 Training Loss: 0.943676 Validation Loss: 0.498315
 Validation loss decreased (0.525996 --> 0.498315). Saving model ...
 Epoch 15, Batch 1 loss: 1.019009
 Epoch 15, Batch 101 loss: 0.908760
 Epoch 15, Batch 201 loss: 0.900598
 Epoch 15, Batch 301 loss: 0.898875
 Epoch: 15 Training Loss: 0.905018 Validation Loss: 0.492271
 Validation loss decreased (0.498315 --> 0.492271). Saving model ...

```

Epoch 16, Batch 1 loss: 1.040815
Epoch 16, Batch 101 loss: 0.893491
Epoch 16, Batch 201 loss: 0.880899
Epoch 16, Batch 301 loss: 0.881799
Epoch: 16          Training Loss: 0.878515          Validation Loss: 0.467434
Validation loss decreased (0.492271 --> 0.467434). Saving model ...
Epoch 17, Batch 1 loss: 1.129126
Epoch 17, Batch 101 loss: 0.870486
Epoch 17, Batch 201 loss: 0.854885
Epoch 17, Batch 301 loss: 0.848426
Epoch: 17          Training Loss: 0.848242          Validation Loss: 0.460464
Validation loss decreased (0.467434 --> 0.460464). Saving model ...
Epoch 18, Batch 1 loss: 0.946332
Epoch 18, Batch 101 loss: 0.871542
Epoch 18, Batch 201 loss: 0.843692
Epoch 18, Batch 301 loss: 0.837989
Epoch: 18          Training Loss: 0.839341          Validation Loss: 0.454109
Validation loss decreased (0.460464 --> 0.454109). Saving model ...
Epoch 19, Batch 1 loss: 1.096290
Epoch 19, Batch 101 loss: 0.810670
Epoch 19, Batch 201 loss: 0.826903
Epoch 19, Batch 301 loss: 0.823902
Epoch: 19          Training Loss: 0.823866          Validation Loss: 0.431243
Validation loss decreased (0.454109 --> 0.431243). Saving model ...
Epoch 20, Batch 1 loss: 0.912636
Epoch 20, Batch 101 loss: 0.822987
Epoch 20, Batch 201 loss: 0.825596
Epoch 20, Batch 301 loss: 0.812630
Epoch: 20          Training Loss: 0.814144          Validation Loss: 0.421388
Validation loss decreased (0.431243 --> 0.421388). Saving model ...

```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [31]: test(loader_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 0.493131
```

```
Test Accuracy: 85% (718/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```

In [32]: from PIL import Image
import torchvision.transforms as transforms

def load_image(img_path):
    image = Image.open(img_path).convert('RGB')
    # resize the input image into 224x224 because VGG16 takes only 224x224 pixel image
    prediction_transform = transforms.Compose([transforms.Resize(size=(224,224)),
                                              transforms.ToTensor()
                                              ])
    image = prediction_transform(image)[:3,:,:].unsqueeze(0)
    return image

In [33]: ### TODO: Write a function that takes a path to an image as input
### and returns the dog breed that is predicted by the model.

# list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in loader_transfer['train'].dataset]

def predict_breed_transfer(model, class_names, img_path):
    # load the image and return the predicted breed
    img = load_image(img_path)
    model = model.cpu()
    model.eval()
    idx = torch.argmax(model(img))
    return class_names[idx]

In [34]: for img_file in os.listdir('./images'):
    img_path = os.path.join('./images', img_file)
    prediction = predict_breed_transfer(model_transfer, class_names, img_path)
    print("image_file_name: {0}, \t prediction breed: {1}".format(img_path, prediction))

image_file_name: ./images/Welsh_springer_spaniel_08203.jpg,          prediction breed: Basset hound
image_file_name: ./images/sample_human_output.png,              prediction breed: Dachshund
image_file_name: ./images/Labrador_retriever_06457.jpg,          prediction breed: Labrador retriever
image_file_name: ./images/Curly-coated_retriever_03896.jpg,       prediction breed: Curly-coated retriever
image_file_name: ./images/sample_cnn.png,                        prediction breed: Nova scotia duck tolling retriever
image_file_name: ./images/Brittany_02625.jpg,                   prediction breed: Brittany
image_file_name: ./images/Labrador_retriever_06449.jpg,          prediction breed: Plott
image_file_name: ./images/American_water_spaniel_00648.jpg,      prediction breed: Cocker spaniel
image_file_name: ./images/sample_dog_output.png,                prediction breed: Great dane
image_file_name: ./images/Labrador_retriever_06455.jpg,          prediction breed: Neapolitan mastiff

```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted



Sample Human Output

breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.


Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm


```
In [35]: ### TODO: Write your algorithm.
### Feel free to use as many code cells as needed.
import matplotlib.pyplot as plt
def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    img = Image.open(img_path)
    plt.imshow(img)
    plt.show()

    if dog_detector(img_path) is True:
        prediction = predict_breed_transfer(model_transfer, class_names, img_path)
        print('Dog is detected !!!!! Ohh i think it looks {0}'.format(prediction))
    elif face_detector(img_path) > 0:
        prediction = predict_breed_transfer(model_transfer, class_names, img_path)
        print("Ohh Human is detected!!!! Human looks like {0}".format(prediction))
    else:
        print("Error!!!! Neither Dogs are found nor Human in the given supplied image")


for img_file in os.listdir('./images'):
    img_path = os.path.join('./images', img_file)
    run_app(img_path)
```




Dog is detected !!!!! Ohh i think it looks Basset hound




Ohh Human is detected!!!!!! Human looks like Dachshund



Dog is detected !!!!!!! Ohh i think it looks Labrador retriever



Dog is detected !!!!! Ohh i think it looks Curly-coated retriever



Error!!!! Neither Dogs are found nor Human in the given supplied image



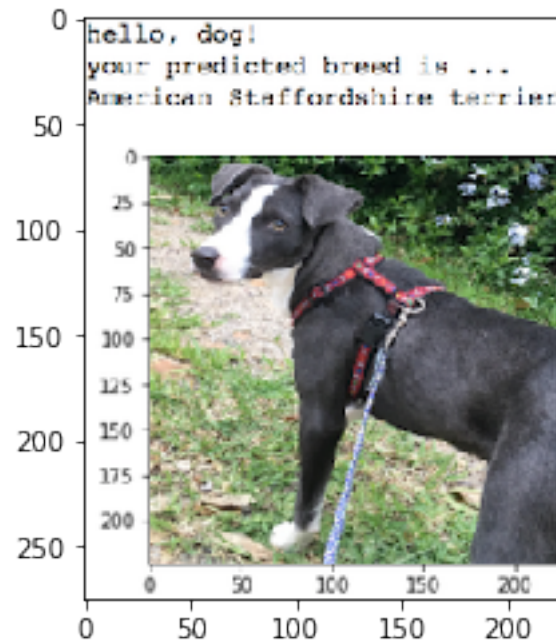
Dog is detected !!!!! Ohh i think it looks Brittany




Dog is detected !!!!! Ohh i think it looks Plott



Dog is detected !!!!! Ohh i think it looks Cocker spaniel



Dog is detected !!!!! Ohh i think it looks Great dane



Dog is detected !!!!! Ohh i think it looks Neapolitan mastiff

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.


Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: I tested my algorithm on atleast 7 images out of which 3 images is of human and 2 images is of cat and 2 is of dog. I got the results as expected. My algorithm accurately predicting the dogs breed as well as predicting the human's breed. For further testing i uploaded 2 images of cats in my dog folder. And my algorithm accurately predicts of that given image. It shows its neither human nor dogs. 1). For improving our model we can do hyperparameters changes to get better and more accurate model. 2). We can do more epochs to get even better accuracy for most complex images of dog's breed.

```
In [36]: human_files = ['./test_images/dp.jpg', './test_images/test_images_1.jpg', './test_images/
dog_files = ['./test_images/dog_1.jpg', './test_images/dog_2.jpg', './test_images/cat_1.j
```

```
In [37]: ## TODO: Execute your algorithm from Step 6 on
## at least 6 images on your computer.
## Feel free to use as many code cells as needed.
```


```
## suggested code, below
for file in np.hstack((human_files[:3], dog_files[:4])):
    run_app(file)
```




Ohh Human is detected!!!! Human looks like American staffordshire terrier




Ohh Human is detected!!!! Human looks like Bull terrier




Ohh Human is detected!!!! Human looks like Bichon frise




Dog is detected !!!!! Ohh i think it looks Black and tan coonhound



Dog is detected !!!!! Ohh i think it looks Plott



Error!!!! Neither Dogs are found nor Human in the given supplied image



Error!!!! Neither Dogs are found nor Human in the given supplied image

In []: